

Deep Learning

Image Classification

Object Detection

Natural Language Modeling

Agenda I

- Introduction Machine Learning
- Neural Networks (NN)
- Overview Deep Learning frameworks
- Introduction PyTorch
- Introduction Convolutional Neural Networks (CNN)
- LeNet
- Introduction ImageNet Challenge
- Various CNN architectures

Agenda II

- Object detection
- Faster R-CNN
- Yolo
- Natural Language Modeling
- Recurrent Neural Networks (RNN)
- Long short-term memory (LSTM)

Introduction Machine Learning I

Supervised Learning	Unsupervised Learning	Reinforcement Learning
Linear Regression Support Vector Machine Decision Tree Neural Network	K-means clustering ...	Q-learning ...

Supervised Learning

1. Goal: find mapping/function from input X to output Y
2. Examples of X/Y pairs are given to train your model

Introduction Machine Learning II




Unknown computation


1. Linear: Linear Regression, Linear Classification
2. Non-linear: Decision Tree, Random Forest, Neural Network

more powerful: map complex X/Y pairs

Introduction Machine Learning III



Linear function



Non-Linear function


Neural Network: Introduction

- Non-linear function approximator
- Can approximate any function arbitrarily close
- Classification and regression


Walter Pitts (1943);
Neural Network pioneer




Neural Network: Building Blocks




Neural Network: Building Blocks




Neural Network: Building Blocks




Neural Network: Building Blocks



Neural Network: Building Blocks




Neural Network: Building Blocks



Neural Network: Universal Approximation Theorem I

“A feed-forward network with a single hidden layer containing a finite number of neurons, can approximate continuous functions on compact subsets of \mathbb{R}^n ”


- Wikipedia

 one hidden layer with lots of neurons can approximate any natural continuous function

Neural Network: Universal Approximation Theorem II

$$y = \sqrt{x_1} + x_2 + e^{x_1} * x_2$$

$$y = \sqrt{x_1 * x_2} + \frac{x_2}{x_1} + x_1 * x_2$$



Neural Network: Universal Approximation Theorem II


A feedforward network with a single layer is sufficient to represent any function, but the layer may be infeasibly large and may fail to learn and generalize correctly.

Ian Goodfellow



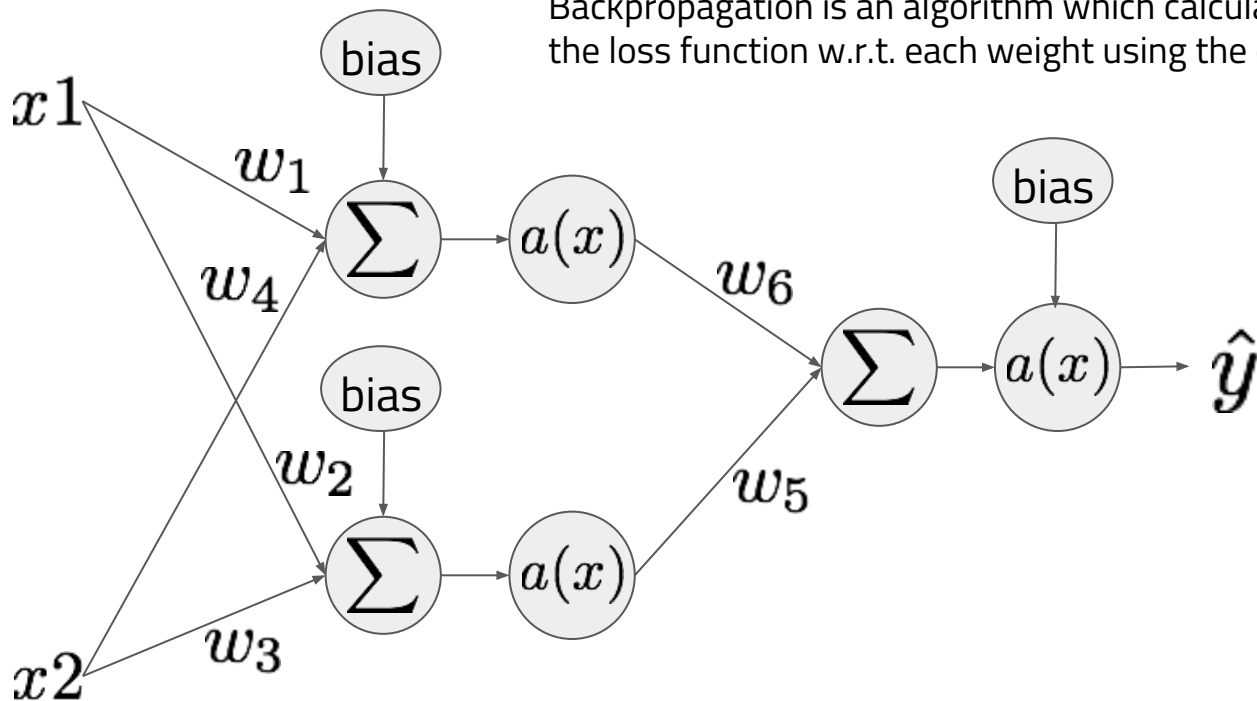
We want our function to generalize well thus we will use more complex architectures containing several hidden layers

Neural Network: How to find ideal weights?



Neural Network: Backpropagation

Backpropagation is an algorithm which calculates the partial derivatives of the loss function w.r.t. each weight using the chain rule.



$$\text{Loss } L = |\hat{y} - y|$$


$$\frac{\partial L}{\partial w_1}$$

...


$$\frac{\partial L}{\partial w_6}$$

Neural Network: Backpropagation

$$\frac{df}{dx} = \frac{df}{dq} * \frac{dq}{dx}$$




Neural Network: Gradient Descent




Neural Network: Different activation functions

- Why do we need non-linear activation functions?
- Hidden Layer:
 - Sigmoid function
 - Rectified Linear Unit (ReLU)
 - Leaky Rectified Linear Unit (Leaky ReLU)
- Output Layer:
 - Softmax function


Neural Network: Sigmoid



Neural Network: ReLu



Neural Network: Leaky ReLu



Overview Deep Learning Frameworks



theano

Typical Features of Deep Learning frameworks

- Compute gradients of the cost function with respect to the weights **automatically**
- Predefined activation functions
- Predefined cost functions
- Predefined optimizers
- Dataloading utilities
- Option to move all computations to the GPU

How do Deep Learning frameworks differ?


- Low level vs high level framework
- Static/dynamic computational graph
- More or less user friendly API
- Help from Community
- Speed of adopting new techniques

Overview Deep Learning Frameworks



theano

Introduction: Deep Learning in the cloud




Introduction to PyTorch

Notebook: `Introduction_to_PyTorch.ipynb`

Exercise I: Linear Regression with PyTorch

Notebook: Exercise_1.ipynb


Introduction Convolutional Neural Networks (CNN)



Humans use forms to classify image

➡ CNNs use the same concept

Introduction Convolutional Neural Networks (CNN)



Part of an image

0	0	0	0	0	0	30
0	0	0	0	50	50	50
0	0	0	20	50	0	0
0	0	0	50	50	0	0
0	0	0	50	50	0	0
0	0	0	50	50	0	0
0	0	0	50	50	0	0


Matrix of part of an image



0	0	0	0	0	30	0
0	0	0	0	30	0	0
0	0	0	30	0	0	0
0	0	0	30	0	0	0
0	0	0	30	0	0	0
0	0	0	30	0	0	0
0	0	0	0	0	0	0

Filter

Introduction Convolutional Neural Networks (CNN)



Part of an image

0	0	0	0	0	0	30
0	0	0	0	50	50	50
0	0	0	20	50	0	0
0	0	0	50	50	0	0
0	0	0	50	50	0	0
0	0	0	50	50	0	0
0	0	0	50	50	0	0

Matrix of part of an image




0	0	0	0	0	30	0
0	0	0	0	30	0	0
0	0	0	30	0	0	0
0	0	0	30	0	0	0
0	0	0	30	0	0	0
0	0	0	30	0	0	0
0	0	0	0	0	0	0

Filter

$$\text{Multiply + Add} = (50 \times 30) + (50 \times 30) + (50 \times 30) + (50 \times 30) + (20 \times 30) + (50 \times 30) = 6600$$

Introduction Convolutional Neural Networks (CNN)



Part of an image

0	0	0	0	0	0	0
0	40	0	0	0	0	0
40	0	40	0	0	0	0
40	20	0	0	0	0	0
0	50	0	0	0	0	0
0	0	50	0	0	0	0
25	25	0	50	0	0	0

Matrix of part of an image

*

0	0	0	0	0	30	0
0	0	0	0	30	0	0
0	0	0	30	0	0	0
0	0	0	30	0	0	0
0	0	0	30	0	0	0
0	0	0	30	0	0	0
0	0	0	0	0	0	0

Filter

Multiply + Add = 0

Introduction Convolutional Neural Networks (CNN)

1 _{x1}	1 _{x0}	1 _{x1}	0	0
0 _{x0}	1 _{x1}	1 _{x0}	1	0
0 _{x1}	0 _{x0}	1 _{x1}	1	1
0	0	1	1	0
0	1	1	0	0

Image

4		


Convolved




Filter operation is performed by a new layer: Convolutional Layer

Convolutional Layer

- Kernel-size (aka filter size): size of the filter matrix (always quadratic)
- Stride (aka step size): step size of filter movement
- Padding: adds pixel around the border of an image (0s most often)



Max Pooling Layer subsamples feature map



Convolutional Layer

- Kernel-size/filter size: 3x3, 5x5, 7x7,...
- Stride: step size of the convolution
- Padding: adds pixel around the border of an image (0s most often)

Calculate output size:

$$\frac{height/width + 2 * padding - filtersize}{stride} + 1$$

Number of weights:


filter size * filter size * # channels + 1 bias

Batch Normalization Layer

- Same idea as input normalization
- Reduces dependency on initialization
- Useful side effect: added regularization
- Two step calculation:
 - Calculate mean over batch and subtract it
 - Calculate standard deviation over batch and divide by it
- two learnable parameters enable the network switch off BN

$$\hat{x} = \frac{x - \mu}{\sigma}$$

$$y = \gamma \hat{x} + \beta$$



Common layers of a CNN

- Conv (convolutional layer): extracts features from input/feature map
- Pool (max pooling layer): subsamples feature map
- FC (fully connected layer): combines features
- ReLu: adds non-linear relationship
- Softmax: output layer for multiclass classification
- BN (batch normalization layer): helps the gradient to flow better

LeNet (1998)

- Input -> Conv -> Sigm -> Pool -> Conv -> Sigm -> Pool -> FC -> Sigm -> Output
- Task: recognize handwritten digits for banks



Yann LeCun, first author of LeNet paper;
now at facebook

Next Notebook: LeNet

Notebook: LeNet.ipynb

History of CNN architectures

LeNet (1998)

ImageNet challenge (2010)

AlexNet (2012)

VGGNet (2014)

Inception (2014)

ResNet (2015)


Xception (2016)

Introduction ImageNet Classification Challenge (ILSVRC)


- started in 2006 with the help of Amazon Mechanical Turk
- 3.2 Mio labelled images; 5,247 categories
- ended in 2017
- > 13 Mio labelled images; > 1000 categories
- **free** to download



Winning architectures became deeper every year



AlexNet (2012)



Architecture of AlexNet

AlexNet (2012) - BEGIN OF DEEP LEARNING

- 60 Mio. parameters, 650,000 neurons
- Top 5 error rate of 15,3% on ImageNet challenge
- 1.2 Mio. trainings images => training took 5-6 days on 2 GPUs (90 Epochs)
- Key contribution: kick off deep learning interest + data augmentation

Next notebook: `AlexNet.ipynb`




Alex Krizhevsky,
now at google

VGGNet (2014)

- Invented by **V**isual **G**eometry **G**roup at Oxford
- VGG16 => 16 layer; VGG19 => 19 layer
- 138 Mio. parameters
- Top 5 error rate of 7,3% on ImageNet challenge
- Training for ImageNet took two to three weeks using four Nvidia Titan Black GPUs
- Easy to understand => good baseline
- Available in most libraries
- Key contribution: Deeper networks perform better

VGGNet (2014)



Architecture of VGG


InceptionNet/GoogLeNet (2014)

- 22 layers/100 layers (depending on how you count)
- Main advantage: only 5 Mio. parameters (no FC layers)
- Top 5 error rate of 6.7% on ImageNet challenge
- Key contribution: Inception module



Christian Szegedy,
now at google

InceptionNet/GoogLeNet (2014)

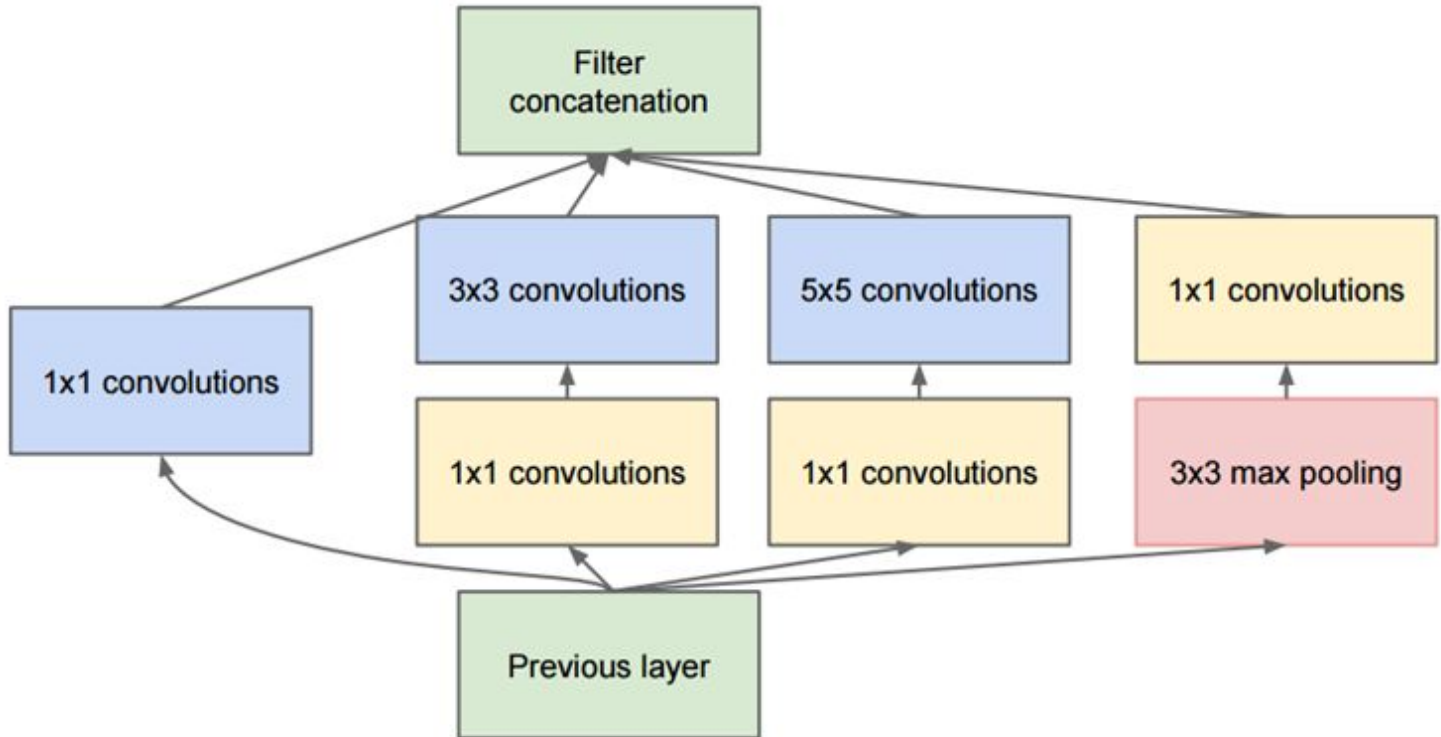


Architecture of InceptionNet

Inception Module

Convolution
Pooling
Softmax
Concat/Normalize

Inception module aka "network in network"



ResNet (2015)

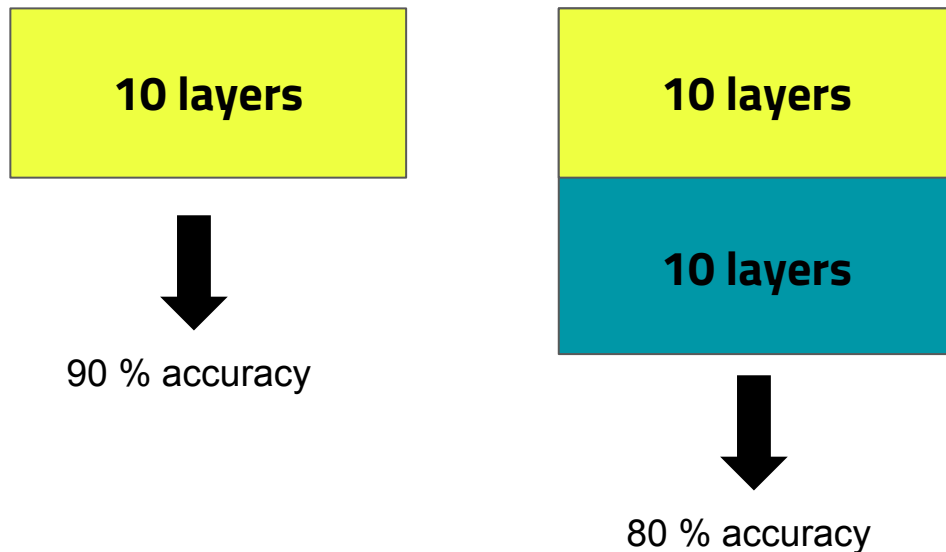
- ResNet50 -> 50 layers; ResNet152 -> 152 layers
- Microsoft Research team
- Top 5 error rate of 3.6% on ImageNet challenge
- Key contribution: Residual block (enables us to make network deeper)



Kaiming He,
now at facebook

ResNet - degradation problem


Thought experiment



This problem is called the degradation problem. Deeper layers struggle to even find an identity mapping.

How can we do deeper?

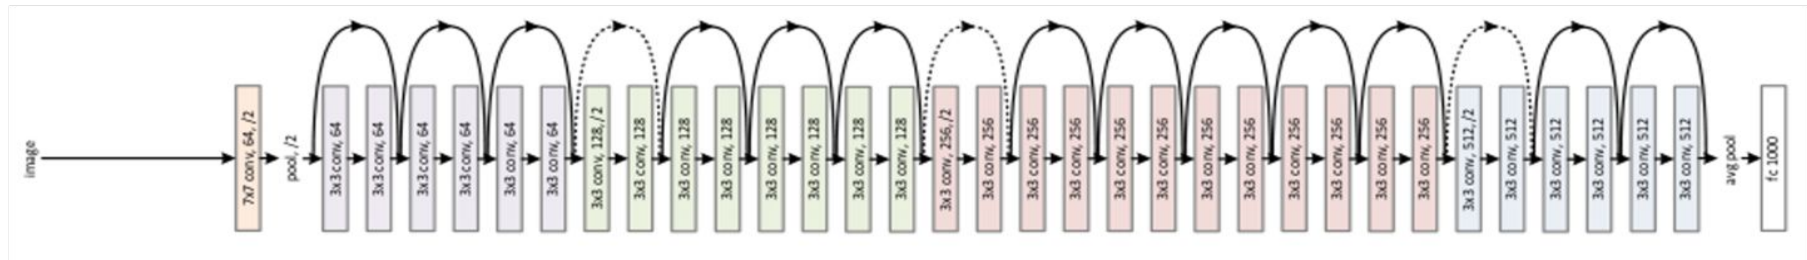
Residual block



Identity is now easy to learn:
 $x = 0 + x$

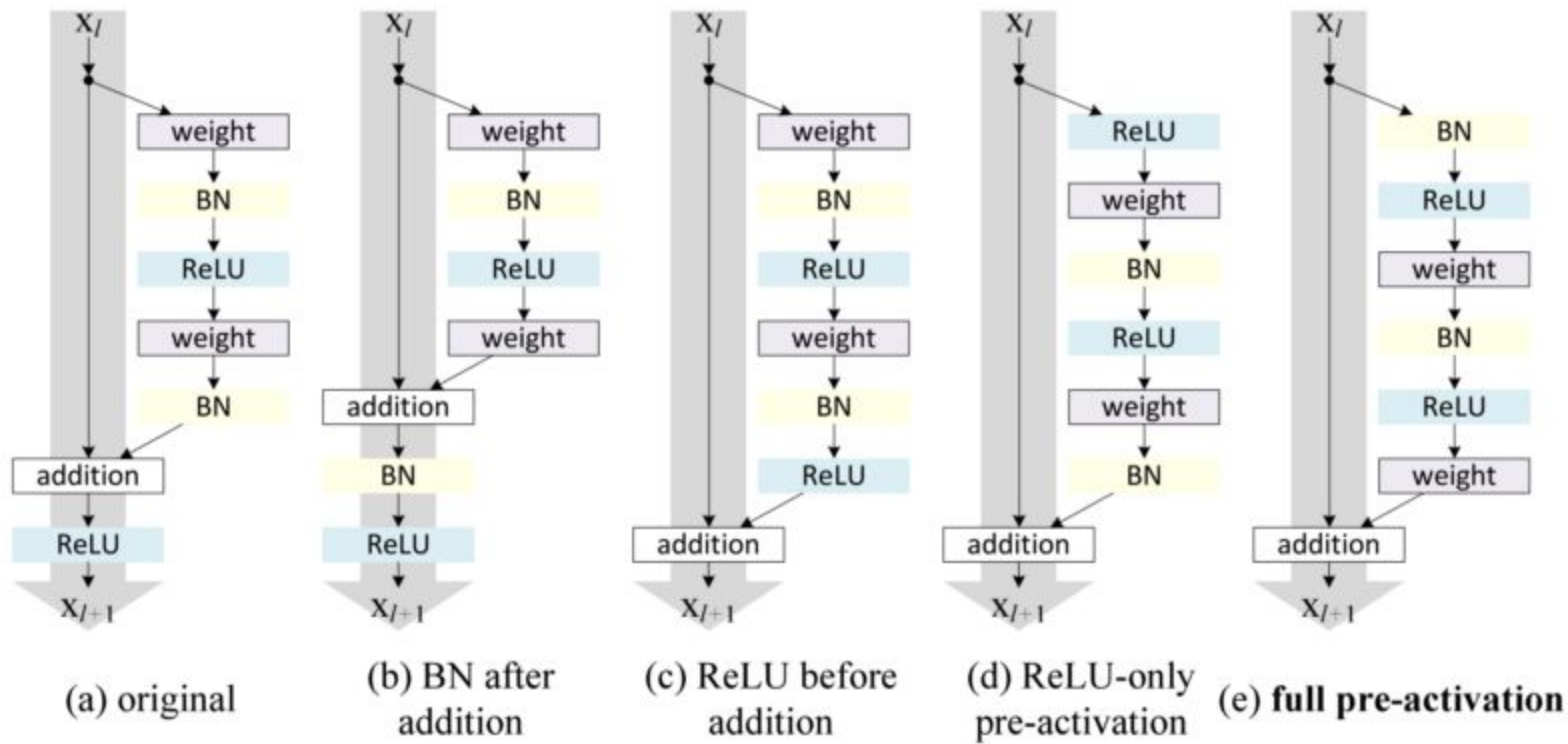
Skip connection

ResNet




Architecture of ResNet

New releases of ResNet




New Hypothesis why ResNet works so well



(a) Conventional 3-block residual network


=



(b) Unraveled view of (a)

➡ ResNet works like an ensemble of many paths

Transfer Learning



➡ Model can now be used on data set B

Next notebook: Transfer Learning


Mittagspause

Introduction Object detection

cat



classification



object detection

mean average precision - mAP

<https://stackoverflow.com/questions/48461855/understanding-and-tracking-of-metrics-in-object-detection>

Regional-based CNN (R-CNN; 2014)

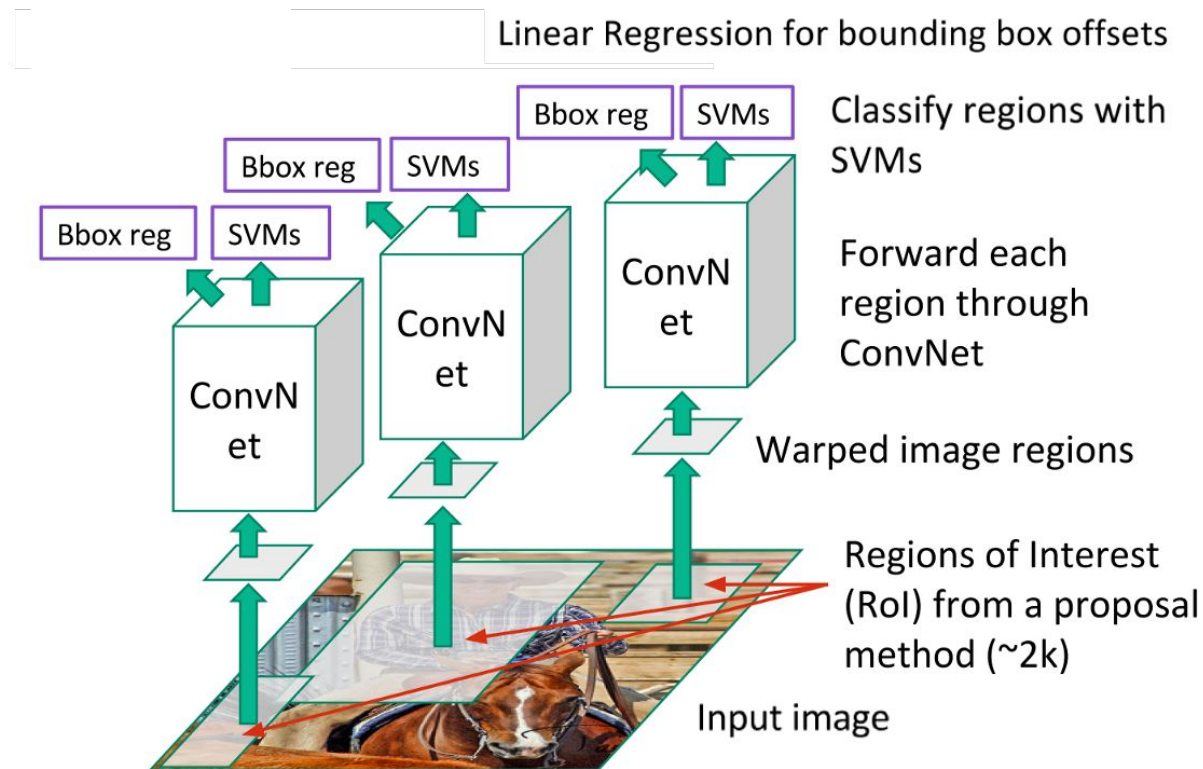
Overview

1. Generate ~ 2000 region proposals by using selective search algorithm
2. Transform each region to a 227x227 image
3. Use CNN to extract a fixed-length feature vector from each region
4. Use SVM + feature vector to classify regions
5. Discard all regions when there is a “close” region with a higher prediction score
6. Adjust size of proposed region with correction factor found by linear classifier

Disadvantage

~47s/image for a single prediction on a GPU  too slow

R-CNN visualized



Fast R-CNN (2015)

Overview

1. Generate regions of interest (ROI) by using selective search algorithm
2. Use one CNN to extract features of the whole input image
3. ROI Pooling: Get extracted features from ROIs
4. Feed ROI features into softmax layer for classification
5. Feed ROI features into another layer and get back four values for the bounding box

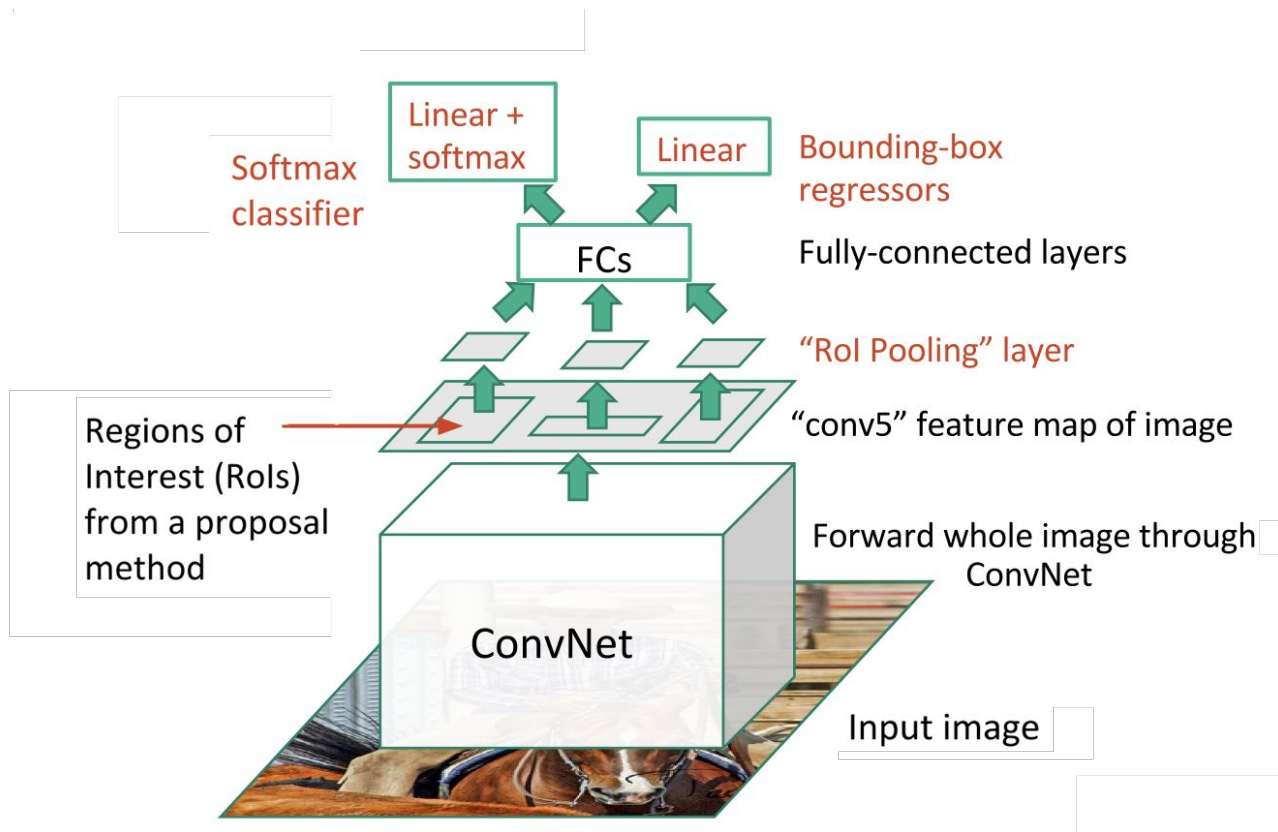
Advantage

All parameters are part of one network



all weights can be adjusted in one step

Fast R-CNN visualized



Faster R-CNN (2015)

Overview

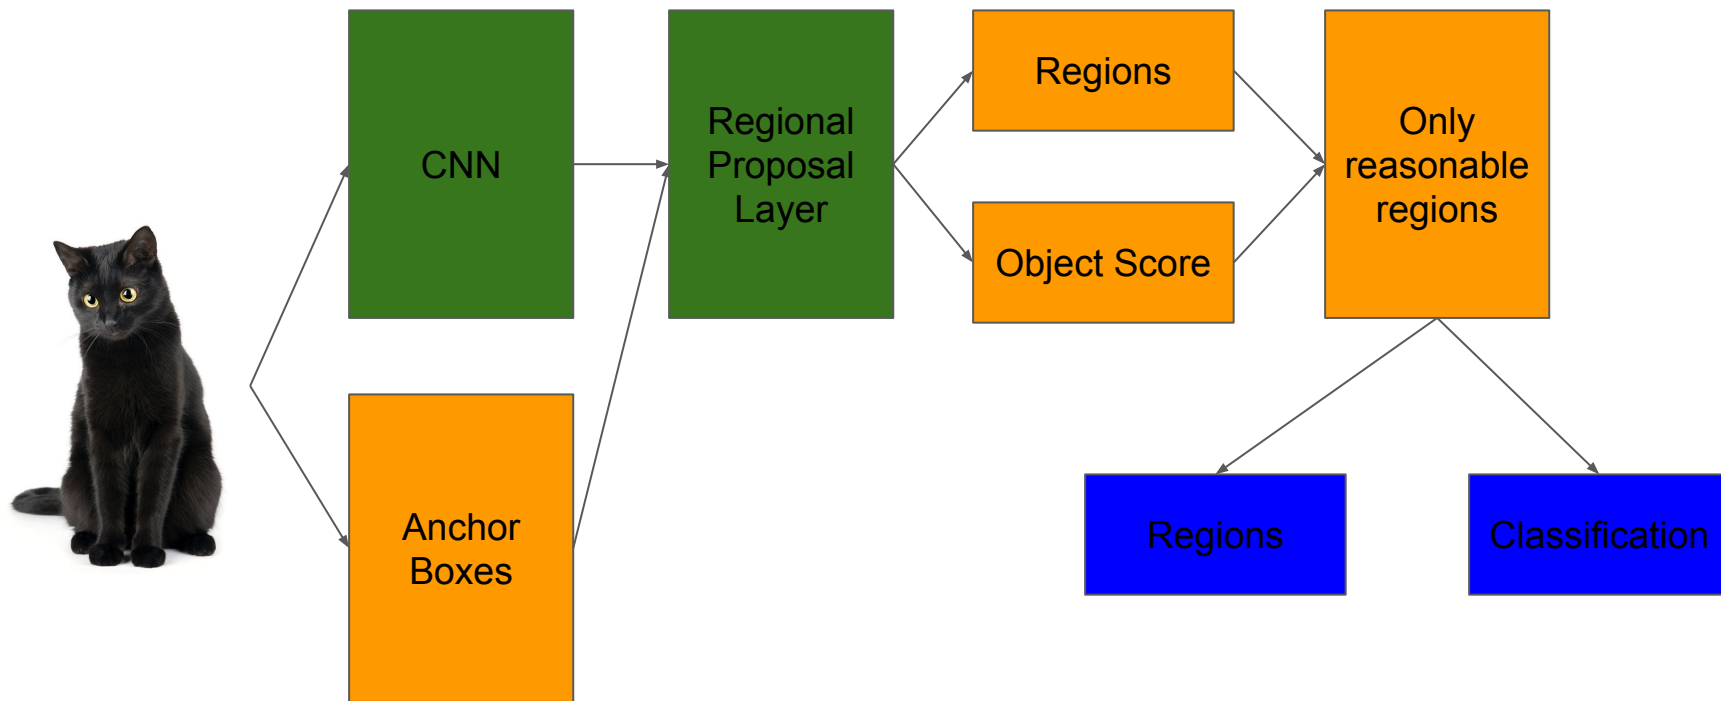
1. Use one CNN to extract features of the whole input image
2. Generate a fixed number of anchor boxes
3. Each anchor box gets a score (object or no object)
4. Depending on the score some anchor boxes will be discarded right away
5. Use the CNN to classify the content of anchor boxes

Advantage

Fully differentiable module


all weights can be adjusted in one step

Faster R-CNN visualized



Faster R-CNN

1. Extract features using a standard CNN




Faster R-CNN


1. Extract features using a standard CNN




2. Every point in the feature map is the center point for k anchor boxes



Faster R-CNN

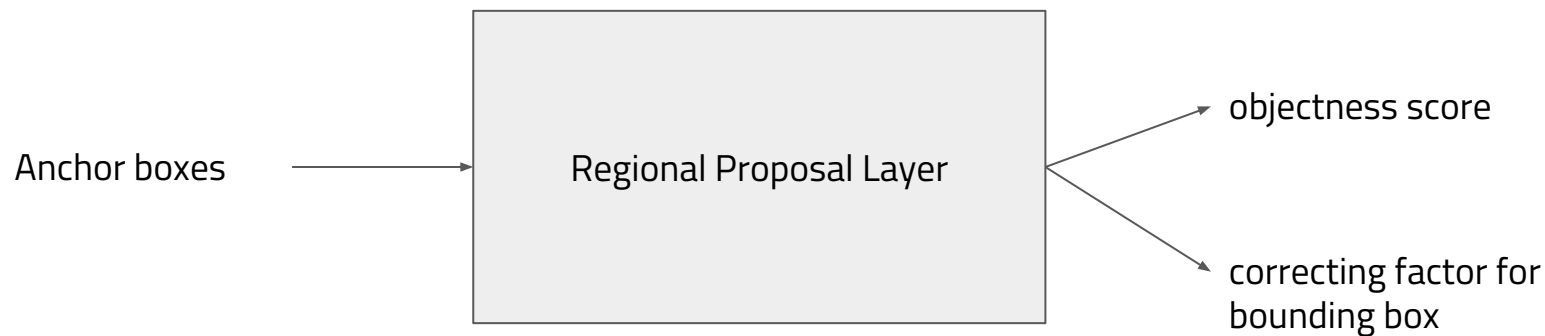


Faster R-CNN




- New center point every 16 pixels
- 9 anchor boxes per center point
- 600 x 800 pic: 17901 anchor boxes

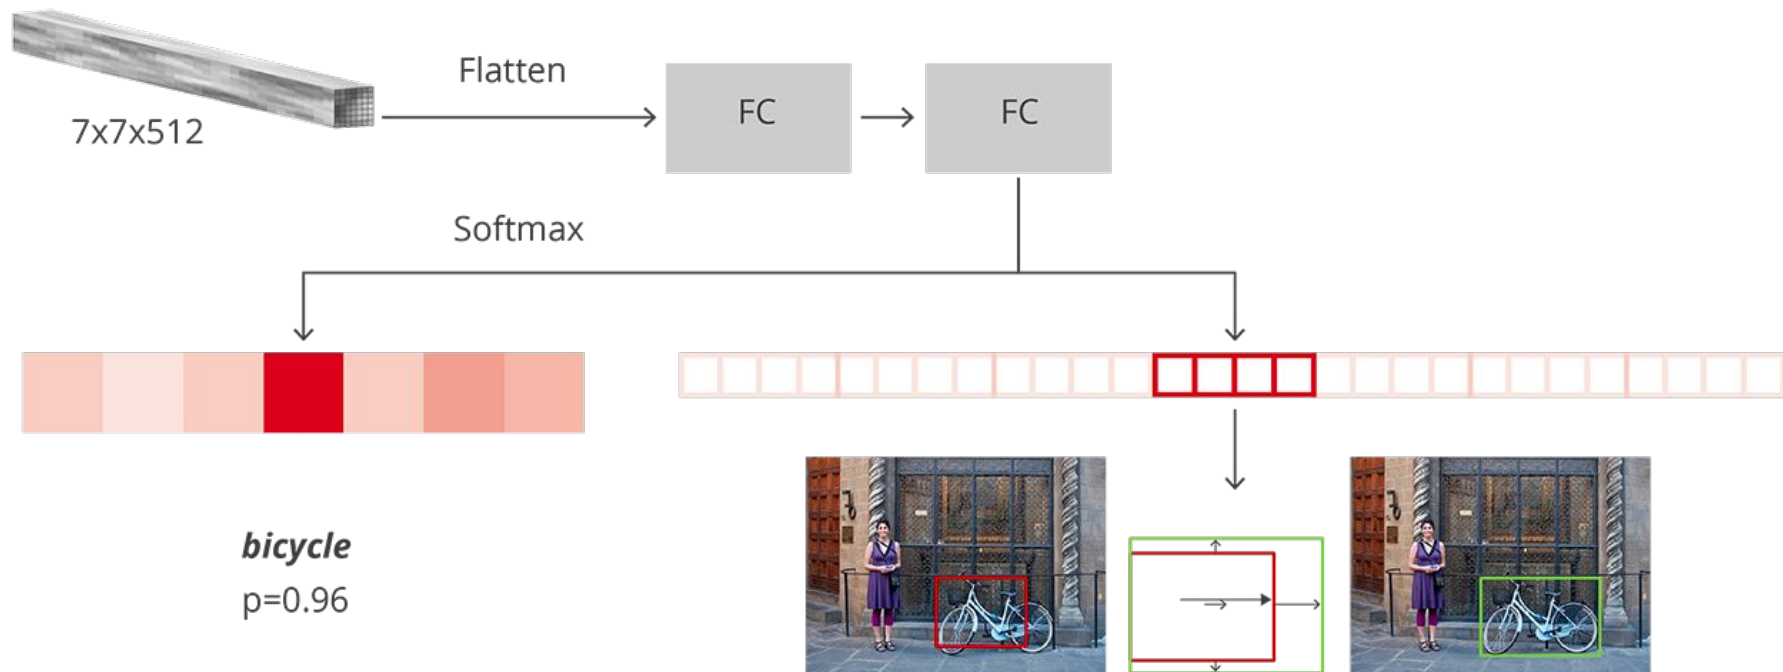
Regional Proposal Network (RPN)



Classify ROI



Final prediction



Training

- Four different losses
 - object score for RPN
 - correction factor for RPN
 - Softmax for final classification
 - correction factor for final bounding box regression
- Researchers found a way to combine these four losses in one loss function

Comparison


R-CNN: 50 secs/image

Fast R-CNN: 2 sec/image

Faster R-CNN: 0.2sec/image

Single Shot Detection

- BEFORE: Region proposal + classification = two steps
- NOW: Regression = one step




Yolo v2


- Divides image in $S \times S$ grid
- Each grid cell predicts
 - B bounding boxes
 - a confidence score for each bounding box
 - One set of class predictions

Bsp: Pascal VOC: $S=7$, $B=2$


$\Rightarrow 7 \times 7 \times (2 \times (4 + 1) + 20 \text{ classes}) = 7 \times 7 \times 30$ output



Yolo v2



Yolo v2



7x7=49 Grid cells, 2 bounding boxes per grid cell

30 infos per grid cell = 2 bounding boxes (x,y,w.h) + 2 confident scores + 20 object classes

Hyperparameter


- Number of grid cells
- Number of bounding boxes per grid cell
- Threshold at which bounding box is considered valid

Comparison Object detection algorithms

Name	backbone	AP_small	AP_medium	AP_large
Faster RCNN	ResNet-101	15.6	38.7	50.9
Yolo v2	Darknet-19	5.0	22.4	35.5
Yolo v3	Darknet-53	18.3	35.4	41.9
RetinaNet	ResNeXt-101	24.1	44.2	51.2

Natural Language Modeling

The Georgetown experiment in 1954 involved fully automatic translation of more than sixty Russian sentences into English. The authors claimed that within three or five years, machine translation would be a solved problem.



Natural Language Modeling - Tasks

- Semantic analysis
- Translation
- Speech2Text
- Text2Speech
- Text generation
- Chat bots
- Spam detection

Word embeddings - Why and What?

- Machine Learning Algorithms can usually only work with numbers not strings
- We could convert every word to a one-hot encoded vector
 - => But relationships between words will be lost

Example

Monday	0	0 0 0 1
Tuesday	1	0 0 1 0
Wednesday	2	0 1 0 0
Thursday	3	1 0 0 0

Word embeddings - Why and What?

- Machine Learning Algorithms usually can only work with numbers not strings
- We could convert every word to a one-hot encoded vector
 - => But relationship between words will be lost

Example

Monday	0	0 0 0 1
Tuesday	1	0 0 1 0
Wednesday	2	0 1 0 0
Thursday	3	1 0 0 0

order remains no relationship

Word embeddings - Why and What?

- Machine Learning Algorithms usually can only work with numbers not strings
- We could convert every word to a one-hot encoded vector
 - => But relationship between words will be lost

Example

Monday
Tuesday
Wednesday
Thursday


0
1
2
3

order remains

0 0 0 1
0 0 1 0
0 1 0 0
1 0 0 0

no relationship

"monkey" 1...000000000000000000000000000000 ...
"flower" 0...000000000001000000000000000000 ...



Pizza



Italy


words as 50000 dimensional vectors

Word embeddings

$$V(\text{"Italy"}) - V(\text{"Pizza"}) = V(\text{"Germany"}) - V(\text{"Bratwurst"})$$

$$V(\text{"King"}) - V(\text{"Queen"}) = V(\text{"Man"}) - V(\text{"Woman"})$$

$$V(\text{"King"}) - V(\text{"Man"}) + V(\text{"Woman"}) = V(\text{"Queen"})$$

 Equations with word vectors actually make sense

Word2Vec

Word2Vec describes a family of algorithms which can generate word embeddings.

- CBOW (Continuous Bag of words)
 - Predict a word given its context words
 - Works well with small amount of the training data, represents well even rare words or phrases
- Skip – Gram model
 - Predict the context words of a given word
 - Several times faster to train than the skip-gram, slightly better accuracy for the frequent words



Tomas Mikolov, currently at Facebook

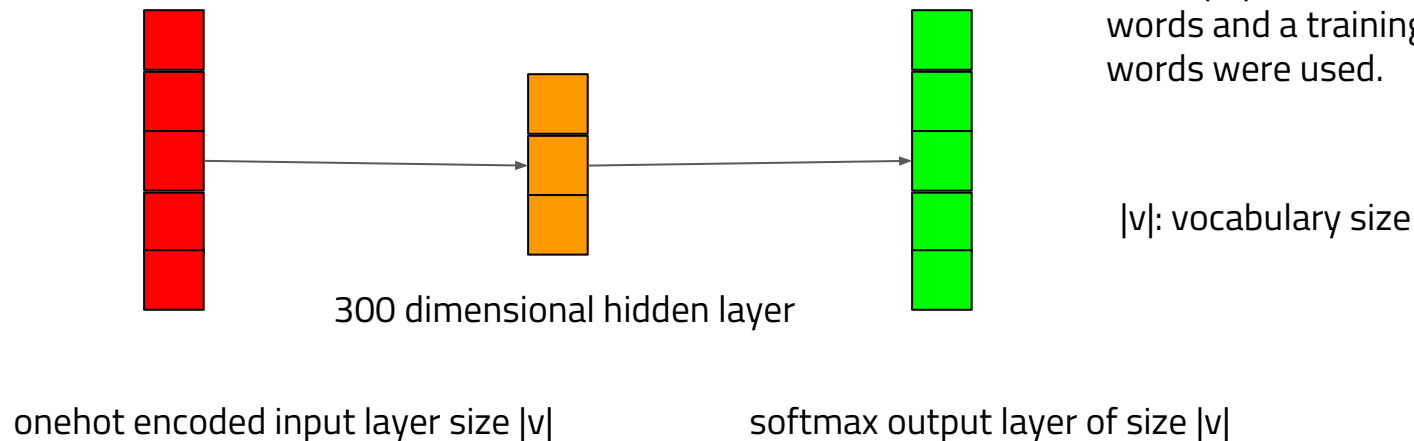


Tomas Mikolov et al managed to speed up the computation of both algorithms tremendously

Word2Vec - Skip gram

Donald Trump really likes watching Fox News and being on the Times cover.

target: watching context: really, liked, Fox News



Global vectors for word representations (GloVe)

GloVe is also an algorithm to build word representations.

1. Build a co-occurrence matrix X
2. Find vectors for the words i and j so that: $\vec{w}_i^T \vec{w}_j + b_i + b_j = \log X_{ij}$

3.
$$J = \sum_{i=1}^V \sum_{j=1}^V f(X_{ij}) (\vec{w}_i^T \vec{w}_j + b_i + b_j - \log X_{ij})^2$$

$$f(X_{ij}) = \begin{cases} \left(\frac{X_{ij}}{x_{\max}} \right)^\alpha & \text{if } X_{ij} < x_{\max} \\ 1 & \text{otherwise.} \end{cases}$$

↖ The function prevents common word pairs to skewing the result

spaCy is a free, open-source library for Natural Language Processing (NLP) in Python.

Main features

- Tokenization
- Lemmatization
- Part-of-Speech tagging
- Entity Recognition
- Word vectors


Next Notebook: Word Embeddings.ipynb



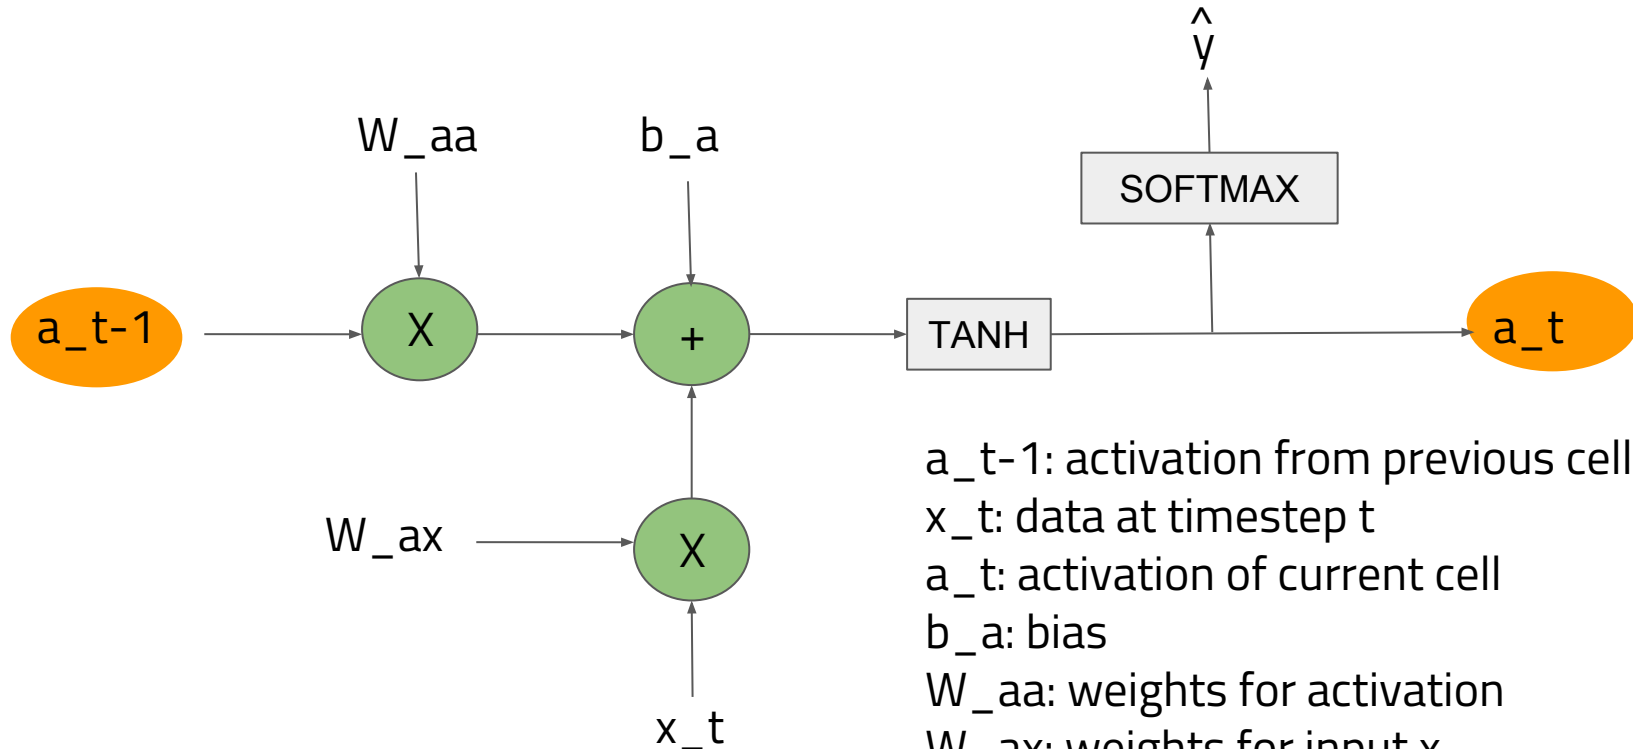
Matt Honnibal, currently at ExplosionAI

Recurrent Neural Network (RNN)

- One RNN consists of several RNN cells
- Useful when data at several timesteps t is available (speech, music, video)




Recurrent Neural Network Cell (RNN)




a_{t-1} : activation from previous cell
 x_t : data at timestep t
 a_t : activation of current cell
 b_a : bias
 W_{aa} : weights for activation
 W_{ax} : weights for input x
 \hat{y} : prediction

RNN - combining several cells


one to one




one to many




many to one



many to many



many to many

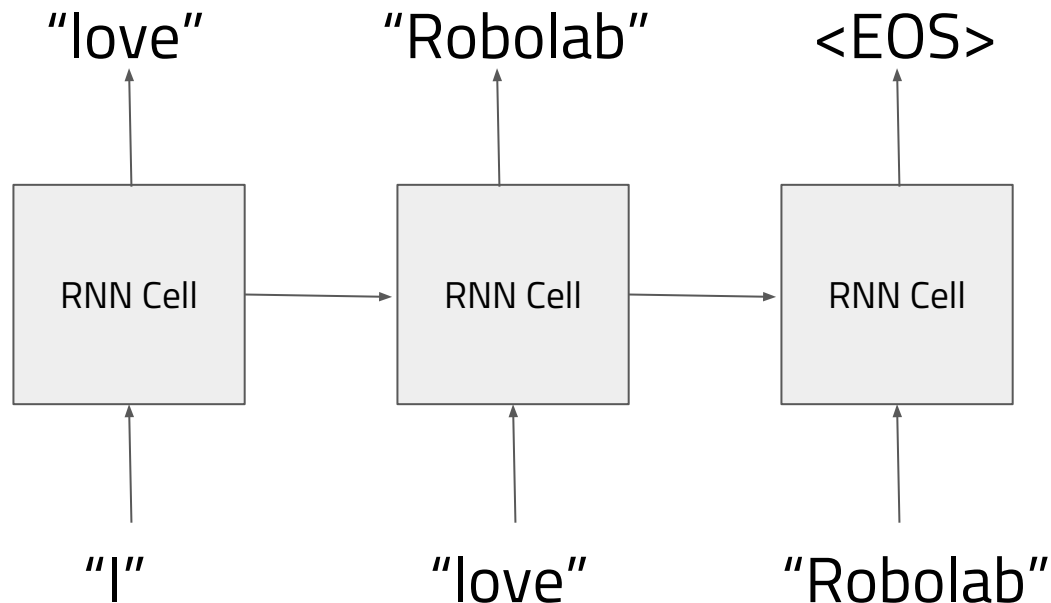


Input layer

RNN cell

Output layer

Generate text with RNNs




Long short-term memory (LSTM)

- RNNs have problems storing a state in a_t over a long period of time
- LSTMs ~ RNNs on steroids
- LSTMs are much more complicated than RNNs
- LSTMs are quite old (1997)

Juergen Schmidhuber, SWISS AI LAB



Long short-term memory (LSTM)



Final Exercise

1. Plant seedling competition (image classification)
2. Rebuild Yolo v2
3. Play around with RNN/LSTMs

